

Basic IR Models

- Inverted Index
- Boolean queries
- Boolean and vector space retrieval model

More realistic scenario

Suppose corpus has 1 million documents(text)

Number of distinct term : 100,000

Now if we create Term Document Incidence matrix for this scenario

$$\text{Number of cell in matrix} = 100,000 * 10,00,000$$

$$= 0.1 * 10^{12}$$

$$= 100\text{GB}$$

(if one cell is stored in 1 byte of memory)



Term-Document Incidence Matrix is Sparse

	Process control block	Process scheduling	CPU utilization	Deadlock in operating system	Disk scheduling algorithm	Critical section
process	1	1	0	0	0	1
kernel	0	1	0	1	0	0
CPU	1	1	1	0	0	0
scheduling	0	1	0	0	1	0
deadlock	0	0	0	1	0	1

The matrix is converted into inverted index

What is an Inverted index ?

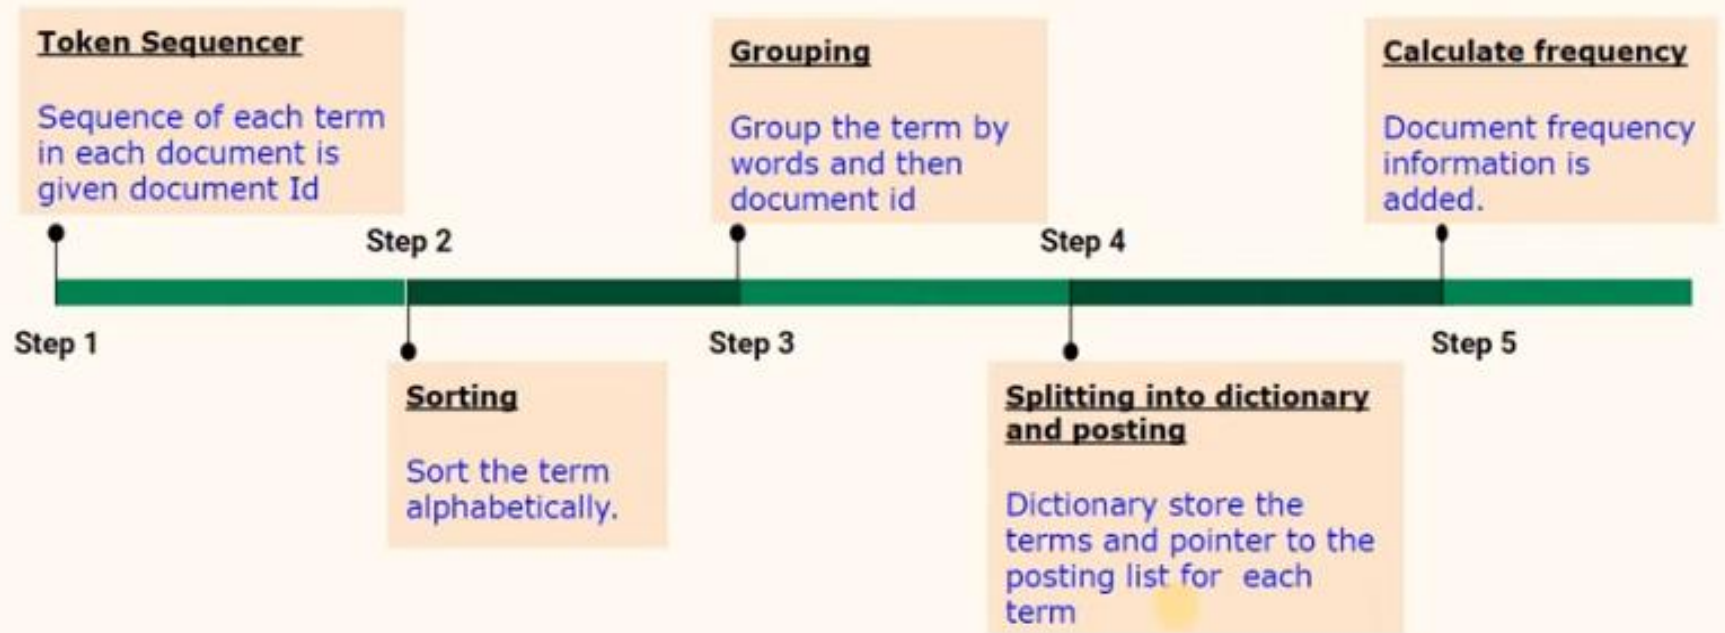
An inverted index is an index data structure storing a mapping from content, such as **words or numbers**, to **its locations in a document or a set of documents**.

It is a hashmap like data structure that directs you **from a word to a document or a web page**.

Inverted Index

Building an Inverted Index

To create an inverted index,



Building an inverted index - Example

Doc1 :

Ram was average student in study but once he became friend of Shyam he started attending lecture regularly.

Doc2:

Shyam was enthusiastic student. He had helped his friend Ram a lot. Now Ram and Shyam are good friends.

Building an inverted index - Example

Term	DocID
Ram	1
was	1
average	1
student	1
in	1
study	1
but	1
once	1
he	1
became	1
friend	1
of	1
Shyam	1
he	1
started	1
attending	1
lecture	1
regularly	1

Term	DocID
Shyam	2
was	2
very	2
enthusiastic	2
student	2
He	2
had	2
helped	2
his	2
friend	2
Ram	2
a	2
lot	2
Now	2
Ram	2
and	2
Shyam	2
are	2
good	2
friend	2



Term	DocID
a	2
and	2
are	2
attending	1
average	1
became	1
but	1
enthusiastic	2
friend	1
friend	2
friend	2
good	2
had	2
he	1
he	1
He	2
helped	2
his	2
in	1
in	1
lecture	1

Term	DocID
lecture	1
lot	2
lot	2
Now	2
Now	2
of	1
once	1
Ram	1
Ram	2
Ram	2
regularly	1
Shyam	1
Shyam	2
Shyam	2
started	1
student	1
student	2
study	1
very	2
was	1
was	2

Building an inverted index - Example

Term	DocID	Term	DocID
a	2	Now	2
and	2	in	1
are	2	lecture	1
attending	1	lot	2
average	1	Now	2
became	1	of	1
but	1	once	1
enthusiastic	2	Ram	1
friend	1	Ram	2
friend	2	Ram	2
friend	2	regularly	1
good	2	Shyam	1
had	2	Shyam	2
he	1	Shyam	2
he	1	started	1
He	2	student	1
helped	2	student	2
his	2	study	1
in	1	very	2
lecture	1	was	1
lot	2	was	2

Term	Doc. Frequency	Posting List
a	1	2
and	1	2
are	1	2
attending	1	1
average	1	1
became	1	1
but	1	1
enthusiastic	1	2
friend	2	1 → 2
good	1	2
had	1	1
he	1	1
He	1	2
helped	1	2
his	1	2
in	1	1
lecture	1	1
lot	1	2
Now	1	2
of	1	1
once	1	1
Ram	2	1 → 2
regularly	1	1
Shyam	2	1
started	1	1
student	2	1
very	1	2
was	2	1

Boolean Retrieval Query

Boolean retrieval model : In this model we represent query which is in boolean expressions of terms.

Terms are combined with **AND** , **NOR** , **NOT**

$X \text{ AND } Y$: represents doc that contains both X and Y .

$X \text{ OR } Y$: represents doc that contains either X or Y .

Processing Boolean queries

Brutus AND Calpurnia

Over the inverted index we:


1. Locate Brutus in the Dictionary
2. Retrieve its postings
3. Locate Calpurnia in the Dictionary
4. Retrieve its postings
5. Intersect the two postings lists, as shown below:

Brutus	→	1	2	4	11	31	45	173	174
Caesar	→	1	2	4	5	6	16	57	132 ...
Calpurnia	→	2	31	54	101				

Implementation(Brutus AND Calpurnia)

INTERSECT(p_1, p_2)

```
1  answer  $\leftarrow \langle \rangle$ 
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
3  do if docID( $p_1$ ) = docID( $p_2$ )
4      then ADD(answer, docID( $p_1$ ))
5           $p_1 \leftarrow \text{next}(p_1)$ 
6           $p_2 \leftarrow \text{next}(p_2)$ 
7  else if docID( $p_1$ ) < docID( $p_2$ )
8      then  $p_1 \leftarrow \text{next}(p_1)$ 
9      else  $p_2 \leftarrow \text{next}(p_2)$ 
10 return answer
```




Answer based on intersection

2

31

Implementation(Brutus OR Calpurnia)



Answer based on Union

1	2	4	11	31	45	54	101	173	174
---	---	---	----	----	----	----	-----	-----	-----

Implementation(NOT Brutus)

Brutus \longrightarrow $\boxed{1} \rightarrow \boxed{2} \rightarrow \boxed{4}$

--	--	--	--	--	--	--	--	--

Answer : based on NOT

Vector Space Model

- The vector space model is a way of representing documents through the words they contain
- It is standard technique in information retrieval
- The VSM allows decisions to be made about which documents are similar to each other and to keyword queries.

Vector Space Model

- In this model, documents and queries are assumed to be part of a t -dimensional vector space, where t is the number of index terms (words, stems, phrases, etc.).
- A document D_i is represented by a vector of index terms:
- $D_i = (d_{i1}, d_{i2}, \dots, d_{it})$,
- where d_{ij} represents the weight of the j th term.

- A document collection containing n documents can be represented as a matrix of term weights, where each row represents a document and each column describes weights that were assigned to a term for a particular document:

	T erm ₁	T erm ₂	...	T erm _{t}
Doc ₁	d_{11}	d_{12}	...	d_{1t}
Doc ₂	d_{21}	d_{22}	...	d_{2t}
.
Doc _{n}	d_{n1}	d_{n2}	...	d_{nt}
	1			

- The **term-document matrix** has been rotated so that now the terms are the rows and the documents are the columns. The term weights are simply the count of the terms in the document.

- D₁ Tropical Freshwater Aquarium Fish.
- D₂ Tropical Fish, Aquarium Care, Tank Setup.
- D₃ Keeping Tropical Fish and Goldfish in Aquariums, and Fish Bowls.
- D₄ The Tropical Tank Homepage - Tropical Fish and Aquariums.

Terms	Documents			
	D ₁	D ₂	D ₃	D ₄
aquarium	1	1	1	1
bowl	0	0	1	0
care	0	1	0	0
fish	1	1	2	1
freshwater	1	0	0	0
goldfish	0	0	1	0
homepage	0	0	0	1
keep	0	0	1	0
setup	0	1	0	0
tank	0	1	0	1
tropical	1	1	1	2

Document D3, for example, is represented by the vector (1, 1, 0, 2, 0, 1, 0, 1, 0, 0, 1).

- **Queries** are represented the same way as documents. That is, a query Q is represented by a vector of t weights:
- $Q = (q_1, q_2, \dots, q_t)$,
- where q_j is the weight of the j th term in the query.
- If, for example the query was “tropical fish”, then using the vector representation in, the query would be $(0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1)$.


- D₁ Tropical Freshwater Aquarium Fish.
- D₂ Tropical Fish, Aquarium Care, Tank Setup.
- D₃ Keeping Tropical Fish and Goldfish in Aquariums, and Fish Bowls.
- D₄ The Tropical Tank Homepage - Tropical Fish and Aquariums.

Terms	Documents			
	D ₁	D ₂	D ₃	D ₄
aquarium	1	1	1	1
bowl	0	0	1	0
care	0	1	0	0
fish	1	1	2	1
freshwater	1	0	0	0
goldfish	0	0	1	0
homepage	0	0	0	1
keep	0	0	1	0
setup	0	1	0	0
tank	0	1	0	1
tropical	1	1	1	2

If, for example the query was “tropical fish”, then using the vector representation in, the query would be (0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1).

- A **similarity measure** is used (rather than a distance or dissimilarity measure), so that the documents with the highest scores are the most similar to the query. The most successful of these is the cosine correlation similarity measure.
- The cosine correlation measures the cosine of the angle between the query and the document vectors.

- Given this representation documents can be ranked by computing the distance between the points representing the documents and query.
 - The documents with the highest score are the most similar to query



Vector Space retrieval Model

- In this model documents and queries are assumed to be part of t dimensional vector space, where t is the no. of index terms (words, phrase etc)
- A document D_i is represented by a vector of index terms:

$$\text{Cosine}(D_i, Q) = \frac{\sum_{j=1}^t d_{ij} \cdot q_j}{\sqrt{\sum_{j=1}^t d_{ij}^2 \cdot \sum_{j=1}^t q_j^2}}$$

- As an example, consider two documents
- $D1 = (0.5, 0.8, 0.3)$ and
- $D2 = (0.9, 0.4, 0.2)$ indexed by three terms, where the numbers represent term weights. Given the query $Q = (1.5, 1.0, 0)$ indexed by the same terms, the cosine measures for the two documents are:

- $\text{Cosine}(D1, Q) = 0.87$
- $\text{Cosine}(D2, Q) = 0.97$
- Since the score is high for document D2
- So we can say document D2 is more relevant to query Q.

TF-IDF(Term frequency-inverse document frequency)

What is TF-IDF(term frequency-inverse document frequency)?

TF-IDF is a statistical measure that evaluates how relevant a word is to a document in a collection of documents.

TF-IDF for a word in a document is calculated by multiplying two different metrics:

- The term frequency of a word in a document(t_f)
- The inverse document frequency(idf) of the word across a set of documents. The closer it is to 0, the more common a word is.

How TF-IDF calculated

- $Tf \cdot idf$ results in TF-IDF score of a word in a document. The higher the score, the more relevant that word in that particular document.
- Mathematical notation:

$$tfidf(t, d, D) = tf(t, d) \cdot idf(t, D)$$

$$tf\ idf(t, d, D) = tf(t, d) \cdot idf(t, D)$$

t=term

d=document

D=set of documents

tf=term frequency

- tf-idf is a weighting scheme that assigns each term in a document a weight based on its term frequency (tf) and inverse document frequency (idf). The terms with higher weight scores are considered to be more important.

- Modern information retrieval by Ricardo Baeza edition 2 page no.100

Vector Space retrieval Model

Inverted index

Tf-ifd

Stemming